

```
#title: "Machine Learning.Rmd"  
#author: "Shivangi"  
#date: "10/14/2021"
```

****Summary****

This document is the final report of the Peer Assessment project from the Practical Machine Learning course, which is a part of the Coursera John's Hopkins University Data Science Specialization. It was written and coded in RStudio, using its knitr functions and published in the html and markdown format. The goal of this project is to predict the manner in which the six participants performed the exercises. The machine learning algorithm, which uses the classe variable in the training set, is applied to the 20 test cases available in the test data.

****Introduction****

Using devices such as Jawbone Up, Nike FuelBand, and Fitbit it is now possible to collect a large amount of data about personal activity relatively inexpensively. These type of devices are part of the quantified self movement – a group of enthusiasts who take measurements about themselves regularly to improve their health, to find patterns in their behavior, or because they are tech geeks. One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it. In this project, your goal will be to use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants. They were asked to perform barbell lifts correctly and incorrectly in 5 different ways.

More information is available from the website here:

![[http://groupware.les.inf.puc-rio.br/har.]]

#Data Source

#The training and test data for this project are collected using the link below:

![[https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv]]

![[https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv]]

#The data for this project come from this source:

![[http://groupware.les.inf.puc-rio.br/har.]]

#The full reference of this data is as follows:

Velloso, E.; Bulling, A.; Gellersen, H.; Ugulino, W.; Fuks, H. "Qualitative Activity Recognition of Weight Lifting Exercises. Proceedings of 4th International Conference in Cooperation with SIGCHI (Augmented Human '13)". Stuttgart, Germany: ACM SIGCHI, 2013.

#Load required R packages and set a seed.

```
```\nlibrary(lattice)\nlibrary(ggplot2)\nlibrary(caret)
```

```
Warning: package 'caret' was built under R version 4.1.1
```

```
library(rpart)\nlibrary(rpart.plot)
```

```
Warning: package 'rpart.plot' was built under R version 4.1.1
```

```
library(corrplot)
```

```
Warning: package 'corrplot' was built under R version 4.1.1
```

```
corrplot 0.90 loaded
```

```
library(rattle)
```

```
Warning: package 'rattle' was built under R version 4.1.1
```

```
Loading required package: tibble
```

```
Warning: package 'tibble' was built under R version 4.1.1
```

```
Loading required package: bitops
```

```
Rattle: A free graphical interface for data science with R.
Version 5.4.0 Copyright (c) 2006-2020 Togaware Pty Ltd.
Type 'rattle()' to shake, rattle, and roll your data.
```

```
library(randomForest)
```

```
Warning: package 'randomForest' was built under R version 4.1.1
```

```
randomForest 4.6-14
```

```
Type rfNews() to see new features/changes/bug fixes.
```

```

Attaching package: 'randomForest'
```

```
The following object is masked from 'package:rattle':

importance
```

```
The following object is masked from 'package:ggplot2':

margin
```

```
library(RColorBrewer)
```

```
Warning: package 'RColorBrewer' was built under R version 4.1.1
```

```
set.seed(222)
```

```
#Load data for training and test datasets.
```

```
url_train <- "http://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv"
url_quiz <- "http://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv"

data_train <- read.csv(url(url_train), strip.white = TRUE, na.strings = c("NA",""))
data_quiz <- read.csv(url(url_quiz), strip.white = TRUE, na.strings = c("NA",""))

dim(data_train)
```

```
[1] 19622 160
```

```
dim(data_quiz)
```

```
[1] 20 160
```

#Create two partitions (75% and 25%) within the original training dataset.

```
in_train <- createDataPartition(data_train$classe, p=0.75, list=FALSE)
train_set <- data_train[in_train,]
test_set <- data_train[-in_train,]

dim(train_set)
```

```
[1] 14718 160
```

```
dim(test_set)
```

```
[1] 4904 160
```

#Two datasets (train\_set and test\_set) have a large number of NA values as well as near-zero-variance (NZV) variables. Both will be removed together with their ID variables.

```
nzv_var <- nearZeroVar(train_set)
train_set <- train_set[, -nzv_var]
test_set <- test_set [, -nzv_var]

dim(train_set)
```

```
[1] 14718 120
```

```
dim(test_set)
```

```
[1] 4904 120
```

#Remove variables that are mostly NA. A threshold of 95 % is selected.

```
na_var <- sapply(train_set, function(x) mean(is.na(x))) > 0.95
train_set <- train_set[, na_var == FALSE]
test_set <- test_set [, na_var == FALSE]

dim(train_set)
```

```
[1] 14718 59
```

```
dim(test_set)
```

```
[1] 4904 59
```

#Since columns 1 to 5 are identification variables only, they will be removed as well.

```
train_set <- train_set[, -(1:5)]
test_set <- test_set [, -(1:5)]

dim(train_set)
```

```
[1] 14718 54
```

```
dim(test_set)
```


```
[1] 4904 54
```

#The number of variables for the analysis has been reduced from the original 160 down to 54.

### **Correlation Analysis**

#Correlation analysis between the variables before the modeling work itself is done. #The “FPC” is used as the first principal component order.

```
corr_matrix <- cor(train_set[, -54])
corrplot(corr_matrix, order = "FPC", method = "circle", type = "lower",
 tl.cex = 0.6, tl.col = rgb(0, 0, 0))
```



#If two variables are highly correlated their colors are either dark blue (for a positive correlation) or dark red (for a negative correlations). Because there are only few strong correlations among the input variables, the Principal Components Analysis (PCA) will not be performed in this analysis. Instead, a few different prediction models will be built to have a better accuracy.

## Prediction Models

### Decision Tree Model

```
set.seed(2222)
fit_decision_tree <- rpart(classe ~ ., data = train_set, method="class")
fancyRpartPlot(fit_decision_tree)
```






```
Confusion Matrix and Statistics
##
Reference
Prediction A B C D E
A 1238 218 37 76 36
B 41 547 28 30 19
C 8 53 688 114 38
D 70 91 50 518 111
E 38 40 52 66 697
##
Overall Statistics
##
Accuracy : 0.752
95% CI : (0.7397, 0.7641)
No Information Rate : 0.2845
P-Value [Acc > NIR] : < 2.2e-16
##
Kappa : 0.685
##
Mcnemar's Test P-Value : < 2.2e-16
##
Statistics by Class:
##
Class: A Class: B Class: C Class: D Class: E
Sensitivity 0.8875 0.5764 0.8047 0.6443 0.7736
Specificity 0.8954 0.9702 0.9474 0.9215 0.9510
Pos Pred Value 0.7713 0.8226 0.7636 0.6167 0.7805
Neg Pred Value 0.9524 0.9052 0.9583 0.9296 0.9491
Prevalence 0.2845 0.1935 0.1743 0.1639 0.1837
Detection Rate 0.2524 0.1115 0.1403 0.1056 0.1421
Detection Prevalence 0.3273 0.1356 0.1837 0.1713 0.1821
Balanced Accuracy 0.8914 0.7733 0.8760 0.7829 0.8623
```

## #Confusion Matrix and Statistics

```
plot(conf_matrix_decision_tree$table, col = conf_matrix_decision_tree$byClass,
 main = paste("Decision Tree Model: Predictive Accuracy =",
 round(conf_matrix_decision_tree$overall['Accuracy'], 4)))
```

## Decision Tree Model: Predictive Accuracy = 0.752



## Generalized Boosted Model (GBM)

```
set.seed(2222)
ctrl_GBM <- trainControl(method = "repeatedcv", number = 5, repeats = 2)
fit_GBM <- train(classe ~ ., data = train_set, method = "gbm",
 trControl = ctrl_GBM, verbose = FALSE)
fit_GBM$finalModel
```

```
A gradient boosted model with multinomial loss function.
150 iterations were performed.
There were 53 predictors of which 53 had non-zero influence.
```

```

predict_GBM <- predict(fit_GBM, newdata = test_set)
conf_matrix_GBM <- confusionMatrix(predict_GBM, factor(test_set$classe))
conf_matrix_GBM

```

```

Confusion Matrix and Statistics
##
Reference
Prediction A B C D E
A 1392 5 0 1 0
B 3 931 4 1 5
C 0 12 843 9 2
D 0 1 8 789 10
E 0 0 0 4 884
##
Overall Statistics
##
Accuracy : 0.9867
95% CI : (0.9831, 0.9898)
No Information Rate : 0.2845
P-Value [Acc > NIR] : < 2.2e-16
##
Kappa : 0.9832
##
Mcnemar's Test P-Value : NA
##
Statistics by Class:
##
Class: A Class: B Class: C Class: D Class: E
Sensitivity 0.9978 0.9810 0.9860 0.9813 0.9811
Specificity 0.9983 0.9967 0.9943 0.9954 0.9990
Pos Pred Value 0.9957 0.9862 0.9734 0.9765 0.9955
Neg Pred Value 0.9991 0.9955 0.9970 0.9963 0.9958
Prevalence 0.2845 0.1935 0.1743 0.1639 0.1837
Detection Rate 0.2838 0.1898 0.1719 0.1609 0.1803
Detection Prevalence 0.2851 0.1925 0.1766 0.1648 0.1811
Balanced Accuracy 0.9981 0.9889 0.9901 0.9884 0.9901

```

## Confusion Matrix and Statistics

## Random Forest Model

```
set.seed(2222)
ctrl_RF <- trainControl(method = "repeatedcv", number = 5, repeats = 2)
fit_RF <- train(classe ~ ., data = train_set, method = "rf",
 trControl = ctrl_RF, verbose = FALSE)
fit_RF$finalModel
```

```
##
Call:
randomForest(x = x, y = y, mtry = min(param$mtry, ncol(x)), verbose = FALSE)
Type of random forest: classification
Number of trees: 500
No. of variables tried at each split: 27
##
OOB estimate of error rate: 0.24%
Confusion matrix:
A B C D E class.error
A 4183 1 0 0 1 0.0004778973
B 8 2836 3 1 0 0.0042134831
C 0 6 2561 0 0 0.0023373588
D 0 0 7 2404 1 0.0033167496
E 0 1 0 7 2698 0.0029563932
```

```
predict_RF <- predict(fit_RF, newdata = test_set)
conf_matrix_RF <- confusionMatrix(predict_RF, factor(test_set$classe))
conf_matrix_RF
```

```

Confusion Matrix and Statistics
##
Reference
Prediction A B C D E
A 1395 3 0 0 0
B 0 946 2 0 0
C 0 0 853 6 0
D 0 0 0 798 1
E 0 0 0 0 900
##
Overall Statistics
##
Accuracy : 0.9976
95% CI : (0.9957, 0.9987)
No Information Rate : 0.2845
P-Value [Acc > NIR] : < 2.2e-16
##
Kappa : 0.9969
##
McNemar's Test P-Value : NA
##
Statistics by Class:
##
Class: A Class: B Class: C Class: D Class: E
Sensitivity 1.0000 0.9968 0.9977 0.9925 0.9989
Specificity 0.9991 0.9995 0.9985 0.9998 1.0000
Pos Pred Value 0.9979 0.9979 0.9930 0.9987 1.0000
Neg Pred Value 1.0000 0.9992 0.9995 0.9985 0.9998
Prevalence 0.2845 0.1935 0.1743 0.1639 0.1837
Detection Rate 0.2845 0.1929 0.1739 0.1627 0.1835
Detection Prevalence 0.2851 0.1933 0.1752 0.1629 0.1835
Balanced Accuracy 0.9996 0.9982 0.9981 0.9961 0.9994

```

## Confusion Matrix and Statistics

### Applying the Best Predictive Model to the Test Data

The following are the predictive accuracy of the three models:

#Decision Tree Model: 75.20 % #Generalized Boosted Model: 98.57 % #Random Forest Model: 99.80 % #The Random Forest model is selected and applied to make predictions on the 20 data points from the original testing dataset (data\_quiz).

```
predict_quiz <- as.data.frame(predict(fit_RF, newdata = data_quiz))
predict_quiz
```

```
predict(fit_RF, newdata = data_quiz)
1 B
2 A
3 B
4 A
5 A
6 E
7 D
8 B
9 A
10 A
11 B
12 C
13 B
14 A
15 E
16 E
17 A
18 B
19 B
20 B
```

```
predict(fit_RF, newdata = data_quiz)
```

```
[1] B A B A A E D B A A B C B A E E A B B B
Levels: A B C D E
```